

# Sean Luka Girgis

Data Engineer | Python, SQL, PySpark, Data Lakes & Pipeline Optimization

214-315-2190 | seanlgirgis@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## Professional Summary

Data Engineer with 20+ years of enterprise technology experience across Python, SQL, PySpark/Spark, data lakes, data warehousing, ETL/ELT pipelines, data modeling, SQL optimization, pipeline tuning, and production analytics systems. Strong background building and improving scalable data workflows, curated reporting datasets, validation checks, performance-oriented schemas, and stakeholder-facing analytics. Best aligned to data engineering roles requiring Python, SQL, PySpark, distributed pipeline processing, data lake/data warehouse delivery, business requirements translation, and cross-functional execution; Azure, Databricks, Data Factory, Power BI, Dataiku, and Azure ML are positioned as adjacent or learning areas rather than overstated production ownership.

## Professional Experience

### Senior Capacity & Data Engineer at CITI (2017-11 – 2025-12)

- Built Python/Pandas/PySpark ETL pipelines ingesting telemetry from 6,000+ infrastructure endpoints and transforming raw operational data into curated datasets for reporting, analytics, and forecasting.
- Designed cloud-oriented data platform components using AWS S3 landing zones, AWS Glue, Athena, Redshift, Oracle, SQL, and PySpark to support scalable transformation, querying, and stakeholder-facing data products.
- Developed large-scale data workflows for ingestion, processing, extraction, validation, and reporting across monitoring, infrastructure, and analytics systems.
- Optimized SQL queries, Oracle schemas, Redshift-backed workloads, indexes, views, historical retention models, and partitioning patterns for performance, maintainability, and cost-aware access.
- Refactored manual and legacy reporting workflows into repeatable ETL processes with documented extraction logic, validation checkpoints, and operational support steps.
- Built data validation, reconciliation, and quality review practices across BMC TrueSight, CA Wily, AppDynamics, Oracle, and cloud-backed reporting layers to improve trust in pipeline outputs.
- Translated business and infrastructure requirements into technical designs, SQL reporting structures, dashboards, and curated data layers for technical and non-technical stakeholders.
- Partnered with infrastructure, monitoring, application, analytics, and business teams to troubleshoot production pipeline issues and deliver reliable data solutions.

### Performance Engineer at G6 Hospitality LLC (2017-03 – 2017-11)

- Developed analytics and reporting outputs from Dynatrace AppMon and Gomez Synthetic Monitoring data for Brand.com and critical hospitality systems.
- Led the FAST project, mining real-user and synthetic performance metrics to identify optimization opportunities and communicate findings to engineering teams.
- Built dashboards and operational reporting views that converted raw monitoring data into actionable performance and reliability insights.

- Supported AWS cloud migration planning by evaluating monitoring readiness, transaction visibility, and data/reporting needs.

#### **SME for CA APM (Senior Consultant) at CA Technologies / TIAA-CREF (2011 – 2016)**

- Managed enterprise telemetry environments with 50+ Enterprise Managers and 4,000–6,000 instrumented agents, supporting reporting, monitoring, and operational analytics.
- Built Perl and Korn shell automation for extracting, transforming, and distributing performance data, replacing manual exports with repeatable reporting workflows.
- Designed dashboards, alert logic, SLA thresholds, technical documentation, and reusable reporting standards adopted across client environments.
- Collaborated with J2EE, WebLogic, infrastructure, database, and operations teams to interpret data, troubleshoot bottlenecks, and improve production visibility.

#### **Performance Test Engineer at AT&T (2010-08 – 2011-07)**

- Analyzed enterprise telecom applications under load to identify throughput limits, SQL/JDBC bottlenecks, CPU, memory, thread, and garbage-collection constraints.
- Configured JMX Monitoring, Wily Introscope, and thread-dump automation to improve runtime visibility during performance and regression testing.
- Documented baseline metrics, design findings, and release-readiness thresholds used by engineering teams for production decisions.

#### **Senior Systems & Data Migration Engineer at Sabre (2008-05 – 2010)**

- Led migration of a high-throughput shopping engine from 200+ MySQL nodes to a 6-node Oracle RAC cluster, preserving transaction performance and data integrity.
- Designed and optimized SQL and Oracle access patterns for high-volume workloads while reducing physical hardware footprint by 95%.
- Built C++/OCCI/OCI validation and regression testing utilities to verify correctness across millions of migrated transaction records.
- Automated schema conversion, data validation, and phased cutover support for a mission-critical enterprise migration.

#### **Architect / Developer (IRS CADE Project) at Computer Science Corporation (CSC) (2007-10 – 2008-05)**

- Designed UML class and sequence diagrams and developed modules for IRS CADE mainframe modernization.
- Built CICS, MQ Series, XML messaging interfaces, and DB2-backed components in a government modernization environment.
- Supported technical design documentation and integration patterns for structured data exchange across modernization components.

#### **Developer / Support Engineer (Billing, Interfaces, CSM/PRMS) at Corpus Inc. / Sprint (2001 – 2007)**

- Developed and supported high-availability C++/Pro\*C/PL/SQL billing interfaces and reporting processes for telecom-scale transaction systems.
- Worked with XML, SQL, Oracle, PL/SQL, messaging, and integration interfaces supporting billing, customer-management, and operational reporting workflows.

- Improved database performance by 10x through SQL, sequence, and process optimization while reducing memory usage by 75%.
- Built automation scripts in Korn shell, Perl, Awk, Syncsort, and SQL to support data processing, reporting, deployments, and operational housekeeping.

## Education

### **Post-Graduate Diploma in Computer Engineering Technology / Computer Science**

Humber College, Toronto, Canada

### **Bachelor of Science in Civil Engineering**

Zagazig University, Egypt

## Skills

**Languages:** Python, SQL, PySpark, PL/SQL, XML, JSON, Perl, Korn Shell, C++, Java

## Flagship Projects

### **Serverless Lakehouse Platform (AWS)**

Designed a scalable AWS data platform pattern using S3, Glue, Athena, Redshift, PySpark, and Parquet for raw ingestion, transformation, curated analytics datasets, query optimization, and reporting access.

**Technologies:** AWS S3, AWS Glue, Athena, Redshift, PySpark, Parquet, SQL

### **HorizonScale — AI Capacity Forecasting Engine**

Built a forecasting-oriented analytics workflow using Python, PySpark, Prophet, scikit-learn, and Streamlit to process infrastructure telemetry, prepare model-ready datasets, and deliver capacity insights.

**Technologies:** Python, PySpark, SQL, Prophet, scikit-learn, Streamlit

### **AI-Powered Job Search Pipeline**

Built a Python automation pipeline with semantic duplicate detection, LLM-based job scoring, JSON artifact generation, quality gates, document rendering, and application tracking using FAISS, sentence-transformers, and LLM APIs.

**Technologies:** Python, FAISS, Sentence Transformers, LLM APIs, JSON, Automation