

Sean Luka Girgis

Senior AWS / PySpark Data Engineer | Glue, Python, SQL & ETL Pipelines

214-315-2190 | seanlgirgis@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Professional Summary

Senior AWS Data Engineer with 20+ years of enterprise technology experience and strong hands-on background in AWS data platforms, Python, SQL, PySpark/Spark, AWS Glue, S3, Athena, Redshift, ETL/ELT pipelines, structured and semi-structured data, data quality, pipeline reliability, performance tuning, and production troubleshooting. Experienced building scalable cloud-oriented data workflows, optimizing SQL and PySpark processing, validating pipeline outputs, and delivering stakeholder-ready analytics in regulated enterprise environments. Best aligned to AWS/PySpark/Glue roles requiring Python, SQL, reusable ETL code, JSON/Parquet data handling, production monitoring, cost-aware design, and reliable pipeline delivery; Terraform is positioned as active learning/working knowledge rather than overstated production ownership.

Professional Experience

Senior Capacity & Data Engineer at CITI (2017-11 – 2025-12)

- Built Python/Pandas/PySpark ETL pipelines ingesting telemetry from 6,000+ infrastructure endpoints and transforming raw operational data into curated datasets for reporting, analytics, forecasting, and validation workflows.
- Designed AWS-based data platform components using S3 landing zones, AWS Glue, Athena, Redshift, Oracle, SQL, and PySpark to support scalable transformation, querying, analytics, and downstream reporting workloads.
- Developed clean, reusable Python, SQL, and PySpark transformation logic for data cleansing, enrichment, aggregation, validation, reporting, and stakeholder-facing analytics.
- Worked with structured and semi-structured operational data using SQL, XML, JSON, Parquet-oriented patterns, schema discipline, partition-aware design, and historical retention models.
- Built data quality, reconciliation, and validation practices across BMC TrueSight, CA Wily, AppDynamics, Oracle, and AWS-backed reporting layers to improve completeness, consistency, and trust in pipeline outputs.
- Optimized SQL queries, Oracle schemas, Redshift-backed workloads, indexes, views, partitioning patterns, and reporting data models for performance, maintainability, and cost-aware processing.
- Monitored, troubleshooted, and improved production data workflows by investigating source feeds, transformation logic, SQL behavior, failed outputs, latency issues, and reporting discrepancies.
- Collaborated with infrastructure, analytics, application, and business stakeholders to translate requirements into reliable data products, dashboards, runbooks, and support documentation.

Performance Engineer at G6 Hospitality LLC (2017-03 – 2017-11)

- Developed analytics and reporting outputs from Dynatrace AppMon and Gomez Synthetic Monitoring data for Brand.com and critical hospitality systems.
- Led the FAST project, mining real-user and synthetic performance metrics to identify optimization opportunities and communicate findings to engineering teams.

- Built dashboards and operational reporting views that converted raw monitoring data into actionable performance and reliability insights.
- Supported AWS cloud migration planning by evaluating monitoring readiness, transaction visibility, and data/reporting needs.

SME for CA APM (Senior Consultant) at CA Technologies / TIAA-CREF (2011 – 2016)

- Managed enterprise telemetry environments with 50+ Enterprise Managers and 4,000–6,000 instrumented agents, supporting monitoring, operational analytics, dashboards, and production reliability.
- Built Perl and Korn shell automation for extracting, transforming, and distributing performance data, replacing manual exports with repeatable reporting workflows.
- Designed dashboards, alert logic, SLA thresholds, technical documentation, and reusable reporting standards adopted across client environments.
- Partnered with J2EE, WebLogic, infrastructure, database, and operations teams to analyze data, troubleshoot bottlenecks, and resolve production issues.

Performance Test Engineer at AT&T (2010-08 – 2011-07)

- Analyzed enterprise telecom applications under load to identify throughput limits, SQL/JDBC bottlenecks, CPU, memory, thread, and garbage-collection constraints.
- Configured JMX Monitoring, Wily Introscope, and thread-dump automation to improve runtime visibility during performance and regression testing.
- Documented baseline metrics, test findings, regression thresholds, and release-readiness notes used by engineering teams for production decisions.

Senior Systems & Data Migration Engineer at Sabre (2008-05 – 2010)

- Led migration of a high-throughput shopping engine from 200+ MySQL nodes to a 6-node Oracle RAC cluster, preserving transaction performance and data integrity.
- Designed and optimized SQL and Oracle access patterns for high-volume workloads while reducing physical hardware footprint by 95%.
- Built C++/OCCI/OCI validation and regression testing utilities to verify correctness across millions of migrated transaction records.
- Automated schema conversion, data validation, and phased cutover support for a mission-critical enterprise migration.

Architect / Developer (IRS CADE Project) at Computer Science Corporation (CSC) (2007-10 – 2008-05)

- Designed UML class and sequence diagrams and developed modules for IRS CADE mainframe modernization.
- Built CICS, MQ Series, XML messaging interfaces, and DB2-backed components in a government modernization environment.
- Supported technical design documentation and integration patterns for structured data exchange across modernization components.

Developer / Support Engineer (Billing, Interfaces, CSM/PRMS) at Corpus Inc. / Sprint (2001 – 2007)

- Developed and supported high-availability C++/Pro*C/PL/SQL billing interfaces and reporting processes for telecom-scale transaction systems.

- Worked with XML, SQL, Oracle, PL/SQL, messaging, and integration interfaces supporting billing, customer-management, and operational reporting workflows.
- Improved database performance by 10x through SQL, sequence, and process optimization while reducing memory usage by 75%.
- Built automation scripts in Korn shell, Perl, Awk, Syncsort, and SQL to support data extraction, transformation, reporting, deployments, and operational housekeeping.

Education

Post-Graduate Diploma in Computer Engineering Technology / Computer Science

Humber College, Toronto, Canada

Bachelor of Science in Civil Engineering

Zagazig University, Egypt

Skills

Languages: Python, SQL, PySpark, PL/SQL, XML, JSON, Perl, Korn Shell, C++, Java

Flagship Projects

Serverless Lakehouse Platform (AWS)

Designed a scalable AWS data platform pattern using S3, Glue, Athena, Redshift, PySpark, and Parquet for raw ingestion, transformation, curated analytics datasets, query optimization, reporting access, and validation checkpoints.

Technologies: AWS S3, AWS Glue, Athena, Redshift, PySpark, Parquet, SQL

HorizonScale — AI Capacity Forecasting Engine

Built a forecasting-oriented analytics workflow using Python, PySpark, Prophet, scikit-learn, and Streamlit to process infrastructure telemetry, prepare model-ready datasets, validate pipeline outputs, and deliver capacity insights.

Technologies: Python, PySpark, SQL, Prophet, scikit-learn, Streamlit, Data Validation

AI-Powered Job Search Pipeline

Built an AI-assisted Python automation pipeline with semantic duplicate detection, LLM-based scoring, JSON artifact generation, quality gates, document rendering, and application tracking using Claude Code, Grok/xAI APIs, FAISS, and sentence-transformers.

Technologies: Python, Claude Code, Grok / xAI API, FAISS, Sentence Transformers, JSON, Quality Gates