

Sean Luka Girgis

Lead SQL / Data Integration Engineer | ETL, Data Quality & Financial Services

214-315-2190 | seanlgirgis@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Professional Summary

Senior Data Engineer and technical delivery lead with 20+ years of enterprise technology experience across SQL, Oracle, PL/SQL, Python, ETL/ELT pipelines, data integrations, data modeling, performance tuning, reconciliation, production support, runbooks, and financial services technology. Strong background building and operating mission-critical data workflows, validating data quality, optimizing database access patterns, documenting operational procedures, and partnering with technical and business stakeholders. Best aligned to lead SQL/data integration roles requiring ETL delivery, data quality controls, auditability, reconciliation, performance tuning, operational maturity, financial-services context, and AI-assisted engineering; SQL Server/T-SQL, Databricks, Snowflake, CDC, and streaming are positioned carefully as adjacent or transferable areas rather than overstated production ownership.

Professional Experience

Senior Capacity & Data Engineer at CITI (2017-11 – 2025-12)

- Built Python/Pandas/PySpark ETL pipelines ingesting telemetry from 6,000+ infrastructure endpoints and transforming raw operational data into curated datasets for reporting, analytics, forecasting, and validation workflows.
- Designed Oracle, SQL, AWS Glue, Athena, Redshift, S3, and PySpark-based data platform components supporting scalable transformation, querying, reporting, and downstream analytics in a regulated financial-services environment.
- Developed SQL-heavy reporting and integration workflows for loading, normalizing, validating, and versioning operational telemetry across monitoring, infrastructure, and analytics data sources.
- Built data quality, reconciliation, and validation practices across BMC TrueSight, CA Wily, AppDynamics, Oracle, and AWS-backed reporting layers to improve completeness, consistency, and trust in pipeline outputs.
- Optimized SQL queries, Oracle schemas, Redshift-backed workloads, indexes, views, partitioning patterns, historical retention models, and reporting data models for performance and maintainability.
- Troubleshoot production data and reporting issues by tracing source feeds, transformation logic, SQL behavior, failed outputs, latency symptoms, and downstream reporting discrepancies.
- Documented transformation rules, metric definitions, data flow notes, runbooks, operational procedures, exception handling patterns, and support steps to improve delivery consistency and knowledge transfer.
- Used AI-assisted development workflows to accelerate Python/SQL work, documentation, quality review, JSON artifact generation, and repeatable delivery practices while preserving engineering review discipline.

Performance Engineer at G6 Hospitality LLC (2017-03 – 2017-11)

- Developed analytics and reporting outputs from Dynatrace AppMon and Gomez Synthetic Monitoring data for Brand.com and critical hospitality systems.

- Led the FAST project, mining real-user and synthetic performance metrics to identify optimization opportunities and communicate findings to engineering teams.
- Built dashboards and operational reporting views that converted raw monitoring data into actionable performance and reliability insights.
- Supported production reliability by correlating application telemetry, transaction behavior, and infrastructure performance trends.

SME for CA APM (Senior Consultant) at CA Technologies / TIAA-CREF (2011 – 2016)

- Managed enterprise telemetry environments with 50+ Enterprise Managers and 4,000–6,000 instrumented agents, supporting monitoring, operational analytics, dashboards, and production reliability.
- Built Perl and Korn shell automation for extracting, transforming, and distributing performance data, replacing manual exports with repeatable reporting workflows.
- Established reusable dashboard, alert, SLA threshold, reporting, and technical documentation standards adopted across client environments.
- Mentored technical teams and partnered with J2EE, WebLogic, infrastructure, database, and operations groups to analyze data, troubleshoot bottlenecks, and resolve production issues.

Performance Test Engineer at AT&T (2010-08 – 2011-07)

- Analyzed enterprise telecom applications under load to identify throughput limits, SQL/JDBC bottlenecks, CPU, memory, thread, and garbage-collection constraints.
- Configured JMX Monitoring, Wily Introscope, and thread-dump automation to improve runtime visibility during performance and regression testing.
- Documented baseline metrics, test findings, regression thresholds, and release-readiness notes used by engineering teams for production decisions.

Senior Systems & Data Migration Engineer at Sabre (2008-05 – 2010)

- Led migration of a high-throughput shopping engine from 200+ MySQL nodes to a 6-node Oracle RAC cluster, preserving transaction performance and data integrity.
- Designed and optimized SQL and Oracle access patterns for high-volume workloads while reducing physical hardware footprint by 95%.
- Built C++/OCCI/OCI validation and regression testing utilities to verify correctness across millions of migrated transaction records.
- Automated schema conversion, data validation, and phased cutover support for a mission-critical enterprise migration.

Education

Post-Graduate Diploma in Computer Engineering Technology / Computer Science

Humber College, Toronto, Canada

Bachelor of Science in Civil Engineering

Zagazig University, Egypt

Skills

Flagship Projects

AI-Powered Job Search Pipeline

Built an AI-assisted Python automation pipeline with semantic duplicate detection, LLM-based scoring, JSON artifact generation, quality gates, document rendering, and application tracking using Claude Code, Grok/xAI APIs, FAISS, and sentence-transformers.

Technologies: Python, Claude Code, Grok / xAI API, FAISS, Sentence Transformers, JSON, Quality Gates

Serverless Lakehouse Platform (AWS)

Designed a scalable AWS data platform pattern using S3, Glue, Athena, Redshift, PySpark, and Parquet for raw ingestion, transformation, curated analytics datasets, query optimization, reporting access, and validation checkpoints.

Technologies: AWS S3, AWS Glue, Athena, Redshift, PySpark, Parquet, SQL

HorizonScale — AI Capacity Forecasting Engine

Built a forecasting-oriented analytics workflow using Python, PySpark, Prophet, scikit-learn, and Streamlit to process infrastructure telemetry, prepare model-ready datasets, validate pipeline outputs, and deliver capacity insights.

Technologies: Python, PySpark, SQL, Prophet, scikit-learn, Streamlit, Data Validation