

Sean Luka Girgis

Enterprise SRE / Observability Engineer | BMC TrueSight, Dynatrace, Python & Linux

214-315-2190 | seanlgirgis@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Professional Summary

Senior enterprise SRE, observability, performance, and platform support professional with 20+ years of experience across BMC TrueSight/TSCO, Dynatrace AppMon, CA APM/Introscope, AppDynamics, Python, Linux/Unix, monitoring, production troubleshooting, workflow automation concepts, CI/CD concepts, cloud platform support, and operational reporting. Strong background supporting large enterprise environments, building telemetry pipelines and dashboards, reducing troubleshooting friction, documenting runbooks, and partnering across infrastructure, application, product, and engineering teams. Best aligned to SRE/observability/platform-support roles requiring enterprise monitoring, production support, Python/shell automation, Linux, cloud concepts, CI/CD awareness, workflow automation, and ServiceNow/BMC Remedy-style ITSM exposure; Ansible Tower, Terraform, Windows administration, .NET, and Tableau are positioned carefully as adjacent or transferable areas rather than overstated production ownership.

Professional Experience

Senior Capacity & Data Engineer at CITI (2017-11 – 2025-12)

- Supported enterprise monitoring and capacity workflows across BMC TrueSight/TSCO, CA Wily/APM, AppDynamics, Oracle, AWS-backed reporting layers, and large-scale infrastructure telemetry sources.
- Built Python/Pandas/PySpark pipelines ingesting telemetry from 6,000+ infrastructure endpoints and transforming raw performance signals into validated dashboards, reports, forecasting datasets, and executive insights.
- Partnered with infrastructure, application, analytics, product, and engineering stakeholders to translate platform signals into practical recommendations for operational stability and service delivery.
- Supported production issue analysis by investigating telemetry feeds, metric behavior, source data quality, failed outputs, latency symptoms, and downstream reporting discrepancies.
- Created KPI-based dashboards and reporting models that converted utilization, performance, and capacity telemetry into actionable views for technical teams and senior stakeholders.
- Developed runbooks, metric definitions, validation checkpoints, reporting logic, support procedures, and documentation to improve repeatability, knowledge sharing, and team onboarding.
- Used Python, SQL, Linux/Unix-style automation patterns, and Git-based workflow discipline to support repeatable operational reporting, data validation, and troubleshooting workflows.
- Applied forecasting workflows with Prophet and scikit-learn to support proactive capacity planning, bottleneck prediction, and risk mitigation across enterprise infrastructure.

Performance Engineer at G6 Hospitality LLC (2017-03 – 2017-11)

- Managed Dynatrace AppMon and Gomez Synthetic Monitoring for Brand.com and critical hospitality systems, supporting application performance visibility and operational readiness.
- Supported monitoring workflows by configuring dashboards, analyzing synthetic monitoring behavior, and correlating user-facing symptoms with application performance signals.

- Led the FAST project, mining real-user and synthetic monitoring metrics to surface optimization opportunities adopted by engineering teams.
- Integrated HP Performance Center with Dynatrace to correlate load-test behavior with APM telemetry and improve issue diagnosis.

SME for CA APM (Senior Consultant) at CA Technologies / TIAA-CREF (2011 – 2016)

- Served as CA APM/Introscope SME for a large financial-services environment with 50+ Enterprise Managers and 4,000–6,000 instrumented agents.
- Led and coordinated CA APM upgrades, monitoring standards, dashboard rollouts, threshold design, alerting practices, and operational reporting adoption across technical teams.
- Designed custom Management Modules, dashboards, alerts, SLA thresholds, and documentation standards used by application, middleware, infrastructure, and operations teams.
- Built Perl and KornShell automation for extracting, validating, transforming, and distributing APM telemetry, replacing manual performance-reporting workflows.
- Provided technical training, onboarding support, troubleshooting guidance, and documentation to improve adoption of monitoring standards and operational playbooks.

Performance Test Engineer at AT&T (2010-08 – 2011-07)

- Analyzed J2EE telecom applications under load to identify throughput limits, JDBC bottlenecks, thread contention, heap pressure, CPU constraints, and garbage-collection symptoms.
- Configured JMX Monitoring, Wily Introscope, and thread-dump automation scripts to improve runtime visibility during performance and regression testing.
- Documented baseline metrics, regression thresholds, test findings, and release-readiness notes used by engineering teams for production decisions.

Senior Systems & Data Migration Engineer at Sabre (2008-05 – 2010)

- Led migration of a high-throughput shopping engine from 200+ MySQL nodes to a 6-node Oracle RAC cluster, preserving transaction performance and data integrity.
- Optimized SQL and Oracle access patterns for high-volume workloads while reducing physical hardware footprint by 95%.
- Built C++/OCCI/OCI validation and regression testing utilities to verify correctness across millions of migrated transaction records.
- Supported performance-focused cutover planning, validation, and troubleshooting for a mission-critical enterprise migration.

Architect / Developer (IRS CADE Project) at Computer Science Corporation (CSC) (2007-10 – 2008-05)

- Designed UML class and sequence diagrams and developed modules for IRS CADE mainframe modernization.
- Built CICS, MQ Series, XML messaging interfaces, and DB2-backed components in a government modernization environment.
- Supported technical design documentation and integration patterns for structured data exchange across modernization components.

Developer / Support Engineer (Billing, Interfaces, CSM/PRMS) at Corpus Inc. / Sprint (2001 – 2007)

- Developed and supported high-availability C++/Pro*C/PL/SQL billing interfaces and data feed processes for telecom-scale transaction systems.
- Built automation scripts in KornShell/ksh, Perl, Awk, Syncsort, and SQL to support data extraction, transformation, reporting, deployments, and operational housekeeping.
- Improved database performance by 10x through SQL, sequence, and process optimization while reducing memory usage by 75%.
- Supported production troubleshooting, release support, and cross-team issue resolution for billing and customer-management systems.

Education

Post-Graduate Diploma in Computer Engineering Technology / Computer Science

Humber College, Toronto, Canada

Bachelor of Science in Civil Engineering

Zagazig University, Egypt

Skills

Flagship Projects

HorizonScale — Capacity Forecasting Engine

Built Python/PySpark forecasting workflows to process infrastructure telemetry, prepare forecasting datasets, validate pipeline outputs, and deliver capacity and bottleneck insights.

Technologies: Python, PySpark, Prophet, scikit-learn, BMC TrueSight, Telemetry

FAST Performance Analytics Project

Mined Dynatrace and Gomez Synthetic Monitoring data to identify optimization opportunities, improve transaction visibility, and communicate actionable findings to engineering teams.

Technologies: Dynatrace AppMon, Gomez Synthetic Monitoring, Performance Analytics, Dashboards, APM

Enterprise APM Reporting Automation

Built Perl and KornShell/ksh automation for extracting, transforming, validating, and distributing performance telemetry across large CA APM enterprise monitoring environments.

Technologies: CA APM, KornShell/ksh, Perl, SQL, Unix/Linux, Monitoring, Automation